


audio-file-classification

December 3, 2024

1 Audio sample file name classification

1.0.1 Contents



1. Data analysis
2. Preprocessing data

3. Choosing performance metrics
4. Multinomial Naive Bayes
5. Random Forest Classifier
6. LSTM (Keras and Tensorflow)
7. Random Forest Classifier - using LSTM model output + numerical features
8. Model validation on unseen data

1.0.2 1. Data analysis

Data columns

- category: sample category (bass, drum, synth, vocal etc.)
- file_path (eg: BGE_13_Mid_Reece_Bass_1.wav)
- file_size (in bytes)
- duration (in miliseconds)
- sample_rate: number of sample-frames per second (in Hz)
- channels: number of audio channels (2 for stereo, 1 for mono)
- frame_count: duration * sample_rate
- sample_type: information size of a single sample-frame

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import swifter
import nltk
nltk.download('punkt_tab')

from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.naive_bayes import MultinomialNB
from sklearn.ensemble import RandomForestClassifier
from sklearn.preprocessing import StandardScaler

import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, LSTM, Dense
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.preprocessing.text import Tokenizer as KerasTokenizer
from tensorflow.keras.metrics import Precision, Recall
from tensorflow.keras.models import load_model

from joblib import dump, load

%matplotlib inline
```

```
[nltk_data] Downloading package punkt_tab to
[nltk_data] /Users/tashvit/nltk_data...
```

[nltk_data] Package punkt_tab is already up-to-date!

Check first 5 rows

```
[2]: audio_data_df = pd.read_csv('audio-file-category-data.csv')
      audio_data_df.head()
```

```
[2]:
```

	category	file_path	file_size	duration	sample_rate	\
0	bass	BGE_13_Mid_Reece_Bass_1.wav	1240832	4.673016	44100	
1	bass	BGE_14_Mid_Reece_Bass_2.wav	1207952	4.548753	44100	
2	bass	BGE_15_808_Sat_Bass_1.wav	1123810	4.233061	44100	
3	bass	BGE_16_808_Sat_Bass_2.wav	656276	2.468413	44100	
4	bass	FJ_09_Bass_01.wav	701182	2.644036	44100	

	channels	frame_count	sample_type
0	2	206080	PCM_24
1	2	200600	PCM_24
2	2	186678	PCM_24
3	2	108857	PCM_24
4	2	116602	PCM_24

Check number of rows, columns

```
[3]: audio_data_df.shape
```

```
[3]: (34110, 8)
```

Check for duplicate data rows

```
[4]: audio_data_df[audio_data_df.duplicated()]
```

```
[4]:
```

	category	file_path	file_size	\
311	bass	DH4_10_Bass_02_C#.wav	74468	
333	bass	SMH_G_Syn_Bass_13.wav	239660	
6159	drum	CLE_Snare_8.wav	94508	
6176	drum	09_HConga_808.wav	25010	
6231	drum	SDAM_Kick_01.wav	139382	
...	
27878	synth-loop	PLPTB_Voices_Lead_Synth_(Wet)_140_C.wav	3636206	
27879	synth-loop	PLPTB_Wheel_Whine_Lead_132_F#.wav	3853064	
27880	synth-loop	PLPTB_Window_Da_Glide_133_E.wav	3823396	
27881	synth-loop	PLPTB_Window_Sea_of_Blue_133_E.wav	3824964	
27882	synth-loop	PLPTB_Window_Sync_Blaster_133_A.wav	3822884	

	duration	sample_rate	channels	frame_count	sample_type
311	0.281270	44100	2	12404	PCM_24
333	0.905578	44100	2	39936	PCM_24
6159	0.357007	44100	2	15744	PCM_24

6176	0.267007	44100	1	11775	PCM_16
6231	0.495238	44100	2	21840	PCM_24
...
27878	13.714263	44100	2	604799	PCM_24
27879	14.545442	44100	2	641454	PCM_24
27880	14.436100	44100	2	636632	PCM_24
27881	14.436100	44100	2	636632	PCM_24
27882	14.436100	44100	2	636632	PCM_24

[510 rows x 8 columns]

Remove duplicates

```
[5]: # Make copy of dataframe
audio_files_df = audio_data_df.copy()

audio_files_df = audio_files_df.drop_duplicates(subset=['file_path'])

audio_files_df[audio_files_df.duplicated()]
```

```
[5]: Empty DataFrame
Columns: [category, file_path, file_size, duration, sample_rate, channels,
frame_count, sample_type]
Index: []
```

Check new shape

```
[6]: audio_files_df.shape
```

```
[6]: (32929, 8)
```

```
[7]: audio_files_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 32929 entries, 0 to 34109
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  -
0   category        32929 non-null  object
1   file_path       32929 non-null  object
2   file_size       32929 non-null  int64
3   duration        32929 non-null  float64
4   sample_rate     32929 non-null  int64
5   channels        32929 non-null  int64
6   frame_count     32929 non-null  int64
7   sample_type     32929 non-null  object
dtypes: float64(1), int64(4), object(3)
memory usage: 2.3+ MB
```

```
[8]: audio_files_df.describe()
```

```
[8]:
```

	file_size	duration	sample_rate	channels	frame_count
count	3.292900e+04	32929.000000	32929.000000	32929.000000	3.292900e+04
mean	2.617031e+06	10.414803	44547.808922	1.973762	4.637037e+05
std	6.131048e+06	24.170615	2815.429229	0.159846	1.072703e+06
min	3.640000e+03	0.000250	22050.000000	1.000000	1.200000e+01
25%	1.243640e+05	0.457256	44100.000000	2.000000	2.024000e+04
50%	3.840440e+05	1.714286	44100.000000	2.000000	7.717500e+04
75%	2.230166e+06	9.155283	44100.000000	2.000000	4.115040e+05
max	1.336231e+08	378.750000	192000.000000	2.000000	1.670288e+07

See the unique categories,

```
[9]: audio_files_df['category'].unique()
```

```
[9]: array(['bass', 'bass-loop', 'drum', 'drum-loop', 'effects',
        'instrumental-loops', 'instruments', 'kits-n-templates',
        'stems-n-midi', 'synth', 'synth-loop', 'vocal'], dtype=object)
```

```
[10]: audio_files_df['category'].nunique()
```

```
[10]: 12
```

Checking the diversity of the 12 category groups,

```
[11]: audio_files_df.groupby('category').count()
```

```
[11]:
```


	file_path	file_size	duration	sample_rate	channels	\
category						
bass	560	560	560	560	560	
bass-loop	598	598	598	598	598	
drum	6445	6445	6445	6445	6445	
drum-loop	2652	2652	2652	2652	2652	
effects	4801	4801	4801	4801	4801	
instrumental-loops	1211	1211	1211	1211	1211	
instruments	1237	1237	1237	1237	1237	
kits-n-templates	2932	2932	2932	2932	2932	
stems-n-midi	1168	1168	1168	1168	1168	
synth	4640	4640	4640	4640	4640	
synth-loop	1226	1226	1226	1226	1226	
vocal	5459	5459	5459	5459	5459	

	frame_count	sample_type
category		
bass	560	560
bass-loop	598	598
drum	6445	6445
drum-loop	2652	2652

effects	4801	4801
instrumental-loops	1211	1211
instruments	1237	1237
kits-n-templates	2932	2932
stems-n-midi	1168	1168
synth	4640	4640
synth-loop	1226	1226
vocal	5459	5459

Create a countplot of files by category

```
[12]: sns.set_style("whitegrid")
sns.countplot(data=audio_files_df, x='category', color='purple')
plt.xticks(rotation=80)
plt.show()
```



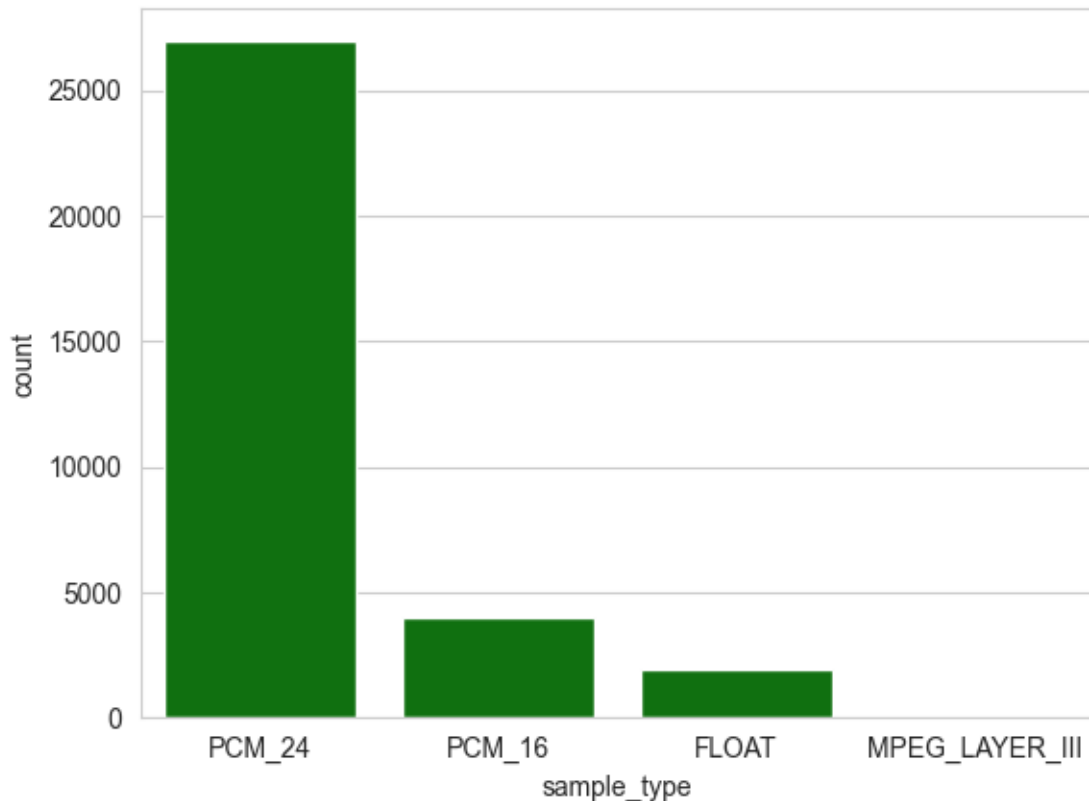
Due to the class imbalance, we may need to oversample minority classes or use class weights.

Checking sample types

```
[13]: audio_files_df['sample_type'].unique()
```

```
[13]: array(['PCM_24', 'PCM_16', 'FLOAT', 'MPEG_LAYER_III'], dtype=object)
```

```
[14]: sns.set_style("whitegrid")
sns.countplot(data=audio_files_df, x='sample_type', color='green')
plt.show()
```



1.0.3 2. Preprocessing data

Convert the category and sample_type columns to numbers

```
[15]: category_mapping = {'bass': 0, 'bass-loop': 1, 'drum': 2, 'drum-loop': 3,
    ↪ 'effects': 4,
    'instrumental-loops': 5, 'instruments': 6, 'kits-n-templates': 7,
    'stems-n-midi': 8, 'synth': 9, 'synth-loop': 10, 'vocal': 11}

sample_type_mapping = {'PCM_24': 0, 'PCM_16': 1, 'FLOAT': 2, 'MPEG_LAYER_III':
    ↪ 3}
```

```

audio_files_df['category_num'] = audio_files_df['category'].
    ↪map(category_mapping)

audio_files_df['sample_type_num'] = audio_files_df['sample_type'].
    ↪map(sample_type_mapping)

audio_files_df.head()

```

```

[15]:
  category      file_path  file_size  duration  sample_rate \
0    bass  BGE_13_Mid_Reece_Bass_1.wav    1240832  4.673016    44100
1    bass  BGE_14_Mid_Reece_Bass_2.wav    1207952  4.548753    44100
2    bass   BGE_15_808_Sat_Bass_1.wav    1123810  4.233061    44100
3    bass   BGE_16_808_Sat_Bass_2.wav     656276  2.468413    44100
4    bass      FJ_09_Bass_01.wav     701182  2.644036    44100

  channels  frame_count sample_type  category_num  sample_type_num
0        2      206080      PCM_24             0                0
1        2      200600      PCM_24             0                0
2        2      186678      PCM_24             0                0
3        2      108857      PCM_24             0                0
4        2      116602      PCM_24             0                0

```

Create a dataframe of file names and categories

```

[16]: audio_text_df = audio_files_df[['file_path', 'category_num']]

audio_text_df.head()

```

```

[16]:
  file_path  category_num
0  BGE_13_Mid_Reece_Bass_1.wav      0
1  BGE_14_Mid_Reece_Bass_2.wav      0
2   BGE_15_808_Sat_Bass_1.wav      0
3   BGE_16_808_Sat_Bass_2.wav      0
4      FJ_09_Bass_01.wav      0

```

Clean the file names

```

[17]: # Make copy of data
audio_filenames_df = audio_files_df.copy()

# Remove the file extension and underscores from file names
def clean_filename(file_name):
    split_filename = file_name.split('.')[0].split('_')
    # if file name end with a number, remove it
    return ' '.join(split_filename[:-1]) if split_filename[-1].isdigit() else ' '
    ↪'.join(split_filename)

```



```
# Using swifter to improve performance
audio_filenames_df['clean_filenames'] = audio_filenames_df['file_path'].swifter.
    ↪allow_dask_on_strings(enable=True).apply(clean_filename)

audio_filenames_df.head(3)
```

Pandas Apply: 0%| | 0/32929 [00:00<?, ?it/s]

```
[17]: category          file_path  file_size  duration  sample_rate \
0      bass  BGE_13_Mid_Reece_Bass_1.wav    1240832  4.673016      44100
1      bass  BGE_14_Mid_Reece_Bass_2.wav    1207952  4.548753      44100
2      bass  BGE_15_808_Sat_Bass_1.wav    1123810  4.233061      44100

      channels  frame_count  sample_type  category_num  sample_type_num \
0           2      206080      PCM_24             0             0
1           2      200600      PCM_24             0             0
2           2      186678      PCM_24             0             0

      clean_filenames
0  BGE 13 Mid Reece Bass
1  BGE 14 Mid Reece Bass
2   BGE 15 808 Sat Bass
```

Tokenise the file names

```
[18]: audio_filenames_df['tokens'] = audio_filenames_df['clean_filenames'].swifter.
    ↪allow_dask_on_strings(enable=True).apply(nltk.word_tokenize)

audio_filenames_df.head(3)
```

Pandas Apply: 0%| | 0/32929 [00:00<?, ?it/s]

```
[18]: category          file_path  file_size  duration  sample_rate \
0      bass  BGE_13_Mid_Reece_Bass_1.wav    1240832  4.673016      44100
1      bass  BGE_14_Mid_Reece_Bass_2.wav    1207952  4.548753      44100
2      bass  BGE_15_808_Sat_Bass_1.wav    1123810  4.233061      44100

      channels  frame_count  sample_type  category_num  sample_type_num \
0           2      206080      PCM_24             0             0
1           2      200600      PCM_24             0             0
2           2      186678      PCM_24             0             0

      clean_filenames          tokens
0  BGE 13 Mid Reece Bass  [BGE, 13, Mid, Reece, Bass]
1  BGE 14 Mid Reece Bass  [BGE, 14, Mid, Reece, Bass]
2   BGE 15 808 Sat Bass  [BGE, 15, 808, Sat, Bass]
```

Extract columns useful for machine learning

```
[19]: # Create dataframe with useful columns for machine learning
audio_data_ml = audio_filenames_df.drop(['category', 'file_path',
↳ 'sample_type', 'clean_filenames'], axis=1)

audio_data_ml.head()
```

```
[19]:
```

	file_size	duration	sample_rate	channels	frame_count	category_num	\
0	1240832	4.673016	44100	2	206080	0	
1	1207952	4.548753	44100	2	200600	0	
2	1123810	4.233061	44100	2	186678	0	
3	656276	2.468413	44100	2	108857	0	
4	701182	2.644036	44100	2	116602	0	

	sample_type_num	tokens
0	0	[BGE, 13, Mid, Reece, Bass]
1	0	[BGE, 14, Mid, Reece, Bass]
2	0	[BGE, 15, 808, Sat, Bass]
3	0	[BGE, 16, 808, Sat, Bass]
4	0	[FJ, 09, Bass]

Remove holdout set to be used for model evaluation

```
[20]: X_all_data = audio_data_ml.drop('category_num', axis=1)
y_all_data = audio_data_ml['category_num']

# Split between seen and unseen data
# (seen data will be used for training and validation, while unseen data can be
↳ used to test the final models)
# Data shuffled automatically
X_seen_data, X_holdout_data, y_seen_data, y_holdout_data =
↳ train_test_split(X_all_data, y_all_data, test_size=0.33, random_state=4567)

X_seen_data.shape, y_seen_data.shape, X_holdout_data.shape, y_holdout_data.shape
```

```
[20]: ((22062, 7), (22062,), (10867, 7), (10867,))
```

Use TfidfVectorizer vectorization to create vocabularies

```
[21]: # Limiting vocabulary size to reduce computational time
MAX_VOCABULARY = 5000
CLASSES = ['bass', 'bass-loop', 'drum', 'drum-loop', 'effects',
↳ 'instrumental-loops', 'instruments', 'kits-n-templates',
↳ 'stems-n-midi', 'synth', 'synth-loop', 'vocal']

def do_nothing(x):
    return x

def baseline_vectorize(documents):
```

```

"""
Create a vectorizer based on given training documents
this is used for the baseline model
:param dataframe: array of word-tokens
:return: vectorizer
"""

# Disable the tokenizer and preprocessor, as it was done in previous
↳pre-processing steps
vectorizer = TfidfVectorizer(tokenizer=do_nothing, preprocessor=do_nothing,
↳lowercase=False, max_features=MAX_VOCABULARY)
# Tokenize and build vocabulary
vectorizer.fit(documents.copy())
# Summarize
print(sorted(vectorizer.vocabulary_)[:50]) # Only show 50 words
print("vocabulary size =", len(vectorizer.vocabulary_))
return vectorizer

```

Building vocabulary from 'X_seen_data'

```

[22]: vectorizer = baseline_vectorize(X_seen_data['tokens'])
X = vectorizer.transform(X_seen_data['tokens'])
y = y_seen_data

```

```

['!', '#', '&', "'ll", "'m", "'re", "'s", '(', ')', '-', '-1', '-2', '-24b',
'-3', '-5', '-6', '0', '00', '000', '0001', '0002', '0004', '0005', '0007',
'0008', '0009', '001', '0010', '0011', '0012', '0014', '0015', '0016', '0017',
'0019', '002', '0020', '0021', '0022', '0023', '0024', '0025', '0026', '0027',
'0028', '0029', '003', '0030', '0031', '0032']
vocabulary size = 5000

```

```

/Users/tashvit/Documents/GitHub/python_fun/.venv/lib/python3.12/site-
packages/sklearn/feature_extraction/text.py:521: UserWarning: The parameter
'token_pattern' will not be used since 'tokenizer' is not None'
warnings.warn(

```

```

[23]: X.shape, y.shape

```

```

[23]: ((22062, 5000), (22062,))

```

Split the seen data into training and testing sets

```

[24]: X_train, X_validation, y_train, y_validation = train_test_split(X, y,
↳test_size=0.2, random_state=6789)

# Check sizes
X_train.shape, X_validation.shape, y_train.shape, y_validation.shape

```

```

[24]: ((17649, 5000), (4413, 5000), (17649,), (4413,))

```

1.0.4 3. Choosing performance metrics

Precision is the most important metric for this multi-class classification problem, as this will minimize false positives.

However a healthy F1-score is also desirable (a high Recall will minimize false negatives).

- True Positives : Instances correctly classified as class .
- False Positives : Instances incorrectly classified as class , but belong to another class.
- False Negatives : Instances that belong to class but were incorrectly classified as another class.

1.1 Analyse performance without oversampling

The author will first analyse performance without oversampling.

If the results are poor, the author will consider oversampling minority classes or use class weights.

1.1.1 4. Multinomial Naive Bayes

```
[25]: mnb = MultinomialNB()  
mnb.fit(X_train.copy(), y_train.copy())
```

```
[25]: MultinomialNB()
```

```
[26]: mnb_predictions = mnb.predict(X_validation)  
  
def print_results(y_true, y_pred):  
    print('Classification report')  
    print(classification_report(y_true, y_pred))  
    print('\n')  
  
    print('Confusion matrix')  
    print(confusion_matrix(y_true, y_pred))  
  
print_results(y_validation, mnb_predictions)
```

Classification report

	precision	recall	f1-score	support
0	0.94	0.27	0.42	60
1	1.00	0.14	0.24	72
2	0.88	0.99	0.93	894
3	0.77	0.87	0.82	354
4	0.88	0.94	0.91	649
5	0.77	0.59	0.67	145
6	0.96	0.50	0.66	161
7	0.83	0.83	0.83	386
8	0.80	0.67	0.73	149

9	0.88	0.99	0.93	637
10	0.84	0.62	0.71	149
11	0.97	0.96	0.97	757
accuracy			0.88	4413
macro avg	0.88	0.70	0.74	4413
weighted avg	0.88	0.88	0.87	4413

Confusion matrix

```
[[ 16  0 19  1  1  1  3  0  0 19  0  0]
 [  0 10  5 23  1  4  0 11  0  7 11  0]
 [  0  0 887  4  3  0  0  0  0  0  0  0]
 [  0  0 29 309  5  0  0  4  5  1  1  0]
 [  0  0 14  8 612  3  0  3  3  1  3  2]
 [  0  0  1  8 15 86  0 16 14  2  2  1]
 [  0  0  8  1  2  1 80  9  0 40  0 20]
 [  0  0 44  8 13  0  0 319  0  2  0  0]
 [  1  0  0 25 11  9  0  0 100  1  1  1]
 [  0  0  0  1  1  0  0  1  0 633  0  1]
 [  0  0  0  8 12  5  0 19  0 13 92  0]
 [  0  0  0  4 16  2  0  1  3  1  0 730]]
```

1.1.2 5. Random Forest Classifier

```
[27]: rfc = RandomForestClassifier(n_estimators=200, n_jobs=-1, random_state=14785,
    ↪ verbose=1)
rfc.fit(X_train.copy(), y_train.copy())
```

```
[Parallel(n_jobs=-1)]: Using backend ThreadingBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 34 tasks      | elapsed:    0.5s
[Parallel(n_jobs=-1)]: Done 184 tasks     | elapsed:    2.3s
[Parallel(n_jobs=-1)]: Done 200 out of 200 | elapsed:    2.5s finished
```

```
[27]: RandomForestClassifier(n_estimators=200, n_jobs=-1, random_state=14785,
    verbose=1)
```

```
[28]: rfc_predictions = rfc.predict(X_validation)

print_results(y_validation, rfc_predictions)
```

Classification report

	precision	recall	f1-score	support
0	0.95	0.90	0.92	60
1	0.94	0.89	0.91	72
2	0.97	0.97	0.97	894
3	0.95	0.94	0.94	354

4	0.97	0.97	0.97	649
5	0.82	0.77	0.79	145
6	0.94	0.72	0.81	161
7	0.80	0.93	0.86	386
8	0.87	0.87	0.87	149
9	0.98	1.00	0.99	637
10	0.91	0.85	0.88	149
11	0.99	0.98	0.99	757
accuracy			0.94	4413
macro avg	0.92	0.90	0.91	4413
weighted avg	0.95	0.94	0.94	4413

Confusion matrix

```
[[ 54  0  0  0  0  0  1  2  0  3  0  0]
 [  3 64  0  2  0  0  0  2  0  0  1  0]
 [  0  0 865  2  1  0  0 26  0  0  0  0]
 [  0  0  7 331  3  1  0  2  5  2  2  1]
 [  0  0  5  1 628  1  0  6  2  1  1  4]
 [  0  2  0  1  6 111  3  4 10  0  7  1]
 [  0  0  1  1  1  0 116 40  0  2  0  0]
 [  0  0 11  6  6  0  0 360  0  3  0  0]
 [  0  0  0  3  0 13  0  3 129  0  0  1]
 [  0  1  0  0  0  0  2  0  0 634  0  0]
 [  0  1  0  2  2  6  1  5  0  5 127  0]
 [  0  0  0  1  3  3  1  0  3  0  1 745]]
```

[Parallel(n_jobs=8)]: Using backend ThreadingBackend with 8 concurrent workers.

[Parallel(n_jobs=8)]: Done 34 tasks | elapsed: 0.0s

[Parallel(n_jobs=8)]: Done 184 tasks | elapsed: 0.0s

[Parallel(n_jobs=8)]: Done 200 out of 200 | elapsed: 0.0s finished

1.1.3 6. LSTM (with Tensorflow and Keras)

Split seen data into training and testing sets using original seen data

```
[29]: X_train, X_validation, y_train, y_validation = train_test_split(X_seen_data,
    ↪ y_seen_data, test_size=0.2, random_state=6789)
```

Function for sequence feature transformation

```
[30]: vocab_size = 500
max_length = 6
embedding_dim = 25

def create_feature_transform(texts):
    """
```

```

Convert text to a sequence that can be used for LSTM training
:param texts: tokens list
:returns: input sequence padded to max_length
"""

tokenizer = KerasTokenizer(num_words=vocab_size, oov_token='<oov>')
tokenizer.fit_on_texts(texts)

def transform(to_transform_text):
    sequences = tokenizer.texts_to_sequences(to_transform_text)
    # Using lstm pad_sequences function to pad the end of the sequence with
    ↪ shorter sequences
    # Longer sequences will be truncated
    return pad_sequences(sequences, maxlen=max_length, padding='post')
return transform

```

Function to train LSTM model

```

[31]: def train_lstm_model(x_features, y_target):
    num_classes = 12 # 0-11, inclusive
    labels = np.array(y_target)
    model = Sequential()
    model.add(Embedding(input_dim=vocab_size, output_dim=embedding_dim))
    model.add(LSTM(units=100))
    model.add(Dense(units=num_classes, activation='softmax')) # Softmax for
    ↪ mutually exclusive classes

    # Compile the model
    model.compile(optimizer='adam',
                  loss='sparse_categorical_crossentropy', # Suitable for
    ↪ integer labels
                  metrics=[Precision(), Recall()])

    # Example training
    model.fit(x_features, labels, epochs=20, batch_size=1)
    return model

# Create a feature transforming function
to_features = create_feature_transform(X_train['tokens'])

# Convert text to features for training and validation
train_features = to_features(X_train['tokens'])
validation_features = to_features(X_validation['tokens'])

```

```
[32]: train_features.shape
```

```
[32]: (17649, 6)
```

Train model using the created function and save it

```
[33]: lstm_model = train_lstm_model(train_features, y_train)
      lstm_model.save('lstm_model.h5') # Save the entire model to a file
```

```
Epoch 1/20
17649/17649          27s 1ms/step
- loss: 0.8420 - precision: 0.9892 - recall: 0.0583
Epoch 2/20
17649/17649          26s 1ms/step
- loss: 0.2731 - precision: 0.9846 - recall: 0.0799
Epoch 3/20
17649/17649          25s 1ms/step
- loss: 0.2020 - precision: 0.9813 - recall: 0.0806
Epoch 4/20
17649/17649          25s 1ms/step
- loss: 0.1512 - precision: 0.9836 - recall: 0.0816
Epoch 5/20
17649/17649          25s 1ms/step
- loss: 0.1340 - precision: 0.9817 - recall: 0.0818
Epoch 6/20
17649/17649          25s 1ms/step
- loss: 0.1256 - precision: 0.9837 - recall: 0.0816
Epoch 7/20
17649/17649          25s 1ms/step
- loss: 0.1062 - precision: 0.9808 - recall: 0.0823
Epoch 8/20
17649/17649          25s 1ms/step
- loss: 0.0866 - precision: 0.9822 - recall: 0.0824
Epoch 9/20
17649/17649          26s 1ms/step
- loss: 0.0797 - precision: 0.9833 - recall: 0.0827
Epoch 10/20
17649/17649          25s 1ms/step
- loss: 0.0711 - precision: 0.9832 - recall: 0.0827
Epoch 11/20
17649/17649          25s 1ms/step
- loss: 0.0729 - precision: 0.9826 - recall: 0.0825
Epoch 12/20
17649/17649          25s 1ms/step
- loss: 0.0699 - precision: 0.9819 - recall: 0.0825
Epoch 13/20
17649/17649          25s 1ms/step
- loss: 0.0650 - precision: 0.9840 - recall: 0.0826
Epoch 14/20
17649/17649          25s 1ms/step
- loss: 0.0625 - precision: 0.9822 - recall: 0.0829
Epoch 15/20
17649/17649          25s 1ms/step
- loss: 0.0643 - precision: 0.9826 - recall: 0.0826
```



```
Epoch 16/20
17649/17649          25s 1ms/step
- loss: 0.0595 - precision: 0.9831 - recall: 0.0828
Epoch 17/20
17649/17649          25s 1ms/step
- loss: 0.0611 - precision: 0.9814 - recall: 0.0827
Epoch 18/20
17649/17649          25s 1ms/step
- loss: 0.0585 - precision: 0.9832 - recall: 0.0829
Epoch 19/20
17649/17649          25s 1ms/step
- loss: 0.0549 - precision: 0.9813 - recall: 0.0828
Epoch 20/20
17649/17649          25s 1ms/step
- loss: 0.0547 - precision: 0.9820 - recall: 0.0829
```

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

```
[34]: # Load the saved model
      lstm_model = load_model('lstm_model.h5')
```

WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to be built. `model.compile_metrics` will be empty until you train or evaluate the model.

Obtain the predicted classes and print results

```
[35]: # Get predicted probabilities
      predicted_probabilities = lstm_model.predict(validation_features)

      # Convert predicted probabilities to class labels
      predicted_classes = predicted_probabilities.argmax(axis=1)

      # Print classification report and confusion matrix
      print_results(y_validation, predicted_classes)
```

```
138/138          0s 1ms/step
Classification report
```

	precision	recall	f1-score	support
0	0.95	0.92	0.93	60
1	0.86	0.86	0.86	72
2	0.98	0.96	0.97	894
3	0.94	0.91	0.92	354
4	0.96	0.96	0.96	649
5	0.85	0.87	0.86	145

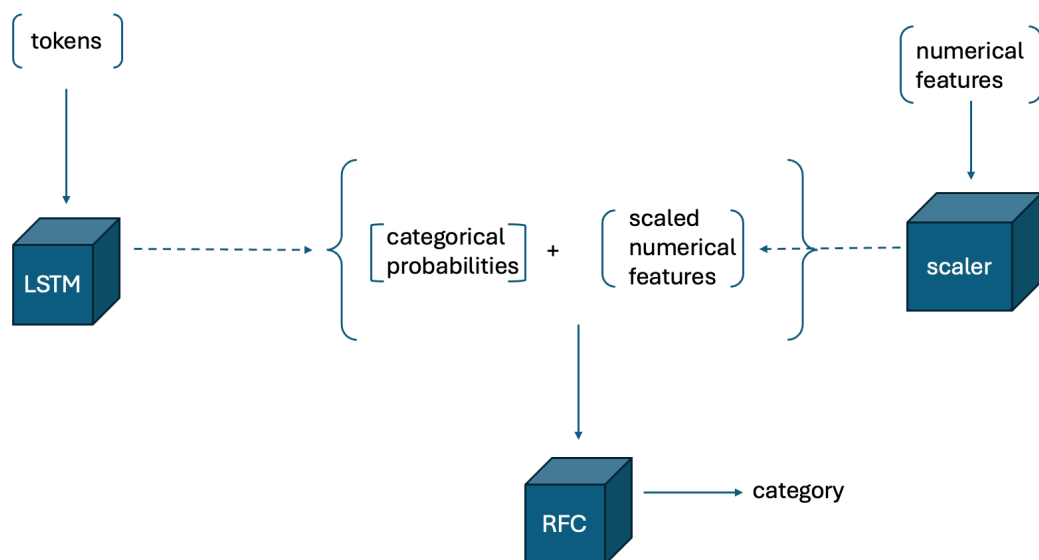
6	0.91	0.89	0.90	161
7	0.85	0.93	0.89	386
8	0.91	0.93	0.92	149
9	0.98	0.98	0.98	637
10	0.86	0.83	0.84	149
11	0.99	0.99	0.99	757
accuracy			0.95	4413
macro avg	0.92	0.92	0.92	4413
weighted avg	0.95	0.95	0.95	4413

Confusion matrix

```
[[ 55   1   1   0   0   0   0   1   0   2   0   0]
 [  0  62   0   3   1   1   0   0   1   0   4   0]
 [  0   0 859   1   4   0   0  29   0   1   0   0]
 [  0   1   9 322   3   2   0   6   7   2   2   0]
 [  0   0   2   3 622   4   0  11   0   1   2   4]
 [  0   3   0   0   5 126   2   0   4   0   4   1]
 [  0   0   0   0   1   2 144   7   0   7   0   0]
 [  0   0   4   3   5   0   9 359   0   1   5   0]
 [  0   0   0   3   0   4   0   4 138   0   0   0]
 [  3   1   0   1   1   0   4   1   0 625   1   0]
 [  0   4   0   8   4   6   0   1   1   1 123   1]
 [  0   0   0   0   3   4   0   2   0   0   2 746]]
```

1.1.4 7. Random Forest Classifier - using LSTM model output + numerical features

This model uses the category probability output from the previous LSTM model with the other numerical features from the dataset.



Getting category probability outputs for 'train_features' derived using file name tokens.

```
[36]: training_probabilities = lstm_model.predict(train_features)

# Print first probability array - 12 probability values for the 12 categories
training_probabilities[0]
```

552/552 0s 658us/step

```
[36]: array([2.1760758e-11, 4.3379949e-09, 9.8595965e-10, 1.9547879e-09,
          3.0972520e-09, 1.9669439e-08, 1.3731193e-08, 8.1804269e-10,
          3.3757136e-10, 1.0000000e+00, 1.6661808e-08, 2.8434726e-09],
          dtype=float32)
```

Pre-process numerical data columns - normalization and standardization

```
[37]: # Check head - 6 numerical features
X_train.head()
```

```
[37]:
```

	file_size	duration	sample_rate	channels	frame_count	\
25864	124194	0.422925	44100	2	18651	
23929	99330	0.328957	44100	2	14507	
21201	4417718	16.695669	44100	2	736279	
19428	16128448	60.952381	44100	2	2688000	
19954	11907338	45.000000	44100	2	1984500	

	sample_type_num	tokens
25864	0	[MPS1, Synth, Shot, 323, B]
23929	0	[MPESS4, Synth, Shot, 088, F]

```

21201          0 [Cymatics, -, Hairspray, -, 115, BPM, G, Maj, ...
19428          0 [WAES, Kit5, Kick, 126, Fm]
19954          0 [WALE2, Kit3, FX, Line02, 128, Gm]

```

```

[38]: numerical_data_X_train = X_train.drop('tokens', axis=1).values
      numerical_data_X_val = X_validation.drop('tokens', axis=1).values

      # Scale the data
      scaler = StandardScaler()

      # Fit the scaler to the training data and transform
      scaled_numerical_data_X_train = scaler.fit_transform(numerical_data_X_train)

      # Save the scaler to a file with joblib
      dump(scaler, 'numerical_scaler.joblib')

```

```

[38]: ['numerical_scaler.joblib']

```

Use the same scaler on the validation data

```

[39]: # Load the saved scaler from the file
      scaler = load('numerical_scaler.joblib')

      # Scale the validation data without fitting
      scaled_numerical_data_X_val = scaler.transform(numerical_data_X_val)

      # Check first entry in scaled training data - 6 values for the 6 numerical
      ↪ features
      scaled_numerical_data_X_train[0]

```

```

[39]: array([-0.40473935, -0.41566295, -0.15263433,  0.16540552, -0.4169775 ,
            -0.43768843])

```

Combine the scaled numerical features with the category probability outputs derived from text features

```

[40]: # Combined features - X train
      X_train_combined_features = np.concatenate([training_probabilities,
      ↪ scaled_numerical_data_X_train], axis=1)

      # Combined features - X validation - use predicted probabilities for X
      ↪ validation
      X_val_combined_features = np.concatenate([predicted_probabilities,
      ↪ scaled_numerical_data_X_val], axis=1)

      # Check first array in combined training feature
      # -> 6 numerical columns + 12 category probabilities = 18 total values
      X_train_combined_features[0]

```

```
[40]: array([ 2.17607581e-11,  4.33799485e-09,  9.85959647e-10,  1.95478789e-09,
           3.09725201e-09,  1.96694394e-08,  1.37311931e-08,  8.18042689e-10,
           3.37571360e-10,  1.00000000e+00,  1.66618079e-08,  2.84347257e-09,
          -4.04739352e-01, -4.15662947e-01, -1.52634330e-01,  1.65405516e-01,
          -4.16977505e-01, -4.37688430e-01])
```

Create a second random forest classifier model object, fit the model and save it.

This random forest classifier model uses the previous LSTM model output and numerical features from the dataset.

```
[41]: # Random forest classifier model that uses previous LSTM model output and
      ↪ numerical features from the dataset
rfc_secondary = RandomForestClassifier(n_estimators=200, n_jobs=-1,
      ↪ random_state=14785, verbose=1)

rfc_secondary.fit(X_train_combined_features.copy(), y_train.copy())

# Save the model to a file with joblib
dump(rfc_secondary, 'rfc_with_lstm_output_num_features.joblib')
```

```
[Parallel(n_jobs=-1)]: Using backend ThreadingBackend with 8 concurrent workers.
```

```
[Parallel(n_jobs=-1)]: Done 34 tasks      | elapsed:    0.1s
```

```
[Parallel(n_jobs=-1)]: Done 184 tasks    | elapsed:    0.6s
```

```
[Parallel(n_jobs=-1)]: Done 200 out of 200 | elapsed:    0.7s finished
```

```
[41]: ['rfc_with_lstm_output_num_features.joblib']
```

Load the saved model

```
[42]: # Load the model from the file
rfc_secondary = load('rfc_with_lstm_output_num_features.joblib')
```

Print classification report and confusion matrix results

```
[43]: rfc_secondary_predictions = rfc_secondary.predict(X_val_combined_features)

print_results(y_validation, rfc_secondary_predictions)
```

Classification report

	precision	recall	f1-score	support
0	0.92	0.93	0.93	60
1	0.89	0.88	0.88	72
2	0.99	0.98	0.98	894
3	0.95	0.93	0.94	354
4	0.96	0.96	0.96	649
5	0.86	0.86	0.86	145
6	0.94	0.92	0.93	161
7	0.92	0.95	0.94	386

8	0.92	0.93	0.93	149
9	0.98	0.98	0.98	637
10	0.84	0.87	0.86	149
11	0.99	0.99	0.99	757
accuracy			0.96	4413
macro avg	0.93	0.93	0.93	4413
weighted avg	0.96	0.96	0.96	4413

Confusion matrix

```
[[ 56  1  1  0  0  0  0  1  0  1  0  0]
 [  0 63  0  2  0  1  0  0  1  0  5  0]
 [  0  0 877  1  5  0  0 10  0  1  0  0]
 [  0  3  3 328  3  0  0  4  7  2  4  0]
 [  0  0  2  3 624  4  0  6  0  1  3  6]
 [  0  0  0  3  5 125  2  0  4  0  5  1]
 [  0  0  0  0  1  1 148  4  0  7  0  0]
 [  0  0  5  3  4  0  1 367  0  1  4  1]
 [  0  0  0  3  0  4  0  3 139  0  0  0]
 [  5  0  0  0  2  0  5  1  0 623  1  0]
 [  0  4  0  3  3  6  0  1  0  1 130  1]
 [  0  0  0  0  3  4  1  1  0  0  2 746]]
```

[Parallel(n_jobs=8)]: Using backend ThreadingBackend with 8 concurrent workers.

[Parallel(n_jobs=8)]: Done 34 tasks | elapsed: 0.0s

[Parallel(n_jobs=8)]: Done 184 tasks | elapsed: 0.0s

[Parallel(n_jobs=8)]: Done 200 out of 200 | elapsed: 0.0s finished

This final attempt has produced the overall best results so far, with a F1 score of 96 and high precision. No data oversampling was necessary.

1.1.5 8. Model validation on unseen data

Validate the final model by testing against the holdout data

```
[44]: # Check head of X holdout data
X_holdout_data.head(3)
```

```
[44]:
```

	file_size	duration	sample_rate	channels	frame_count	\
26026	74770	0.235828	44100	2	10400	
30730	105884	0.400000	44100	2	17640	
31499	137460	0.473061	44100	2	20862	

	sample_type_num	tokens
26026	0	[MPS2, Synth, 117, F, #]
30730	0	[mp1vs2, vocal, shot, 0612, d, #]
31499	0	[MPVGW3, Vocal, Glitch, One, Shot, 081]

Process the unseen data so it can be used by the LSTM model

```
[45]: # Holdout data tokens - convert text to features
holdout_features = to_features(X_holdout_data['tokens'])

holdout_features.shape
```

```
[45]: (10867, 6)
```

Get predicted probabilities from the trained LSTM model, and combine with the scaled numerical features of the unseen data

```
[46]: # Arrays that each contain 12 probability values for the 12 categories
holdout_probabilities = lstm_model.predict(holdout_features)

# Scale the holdout data using the existing Scaler without fitting
numerical_data_X_holdout = X_holdout_data.drop('tokens', axis=1).values
scaled_numerical_data_X_holdout = scaler.transform(numerical_data_X_holdout)

# Combine the scaled numerical features with the category probability outputs
X_holdout_combined_features = np.concatenate([holdout_probabilities,
↪ scaled_numerical_data_X_holdout], axis=1)

# Check first array in combined holdout feature
# -> 6 numerical columns + 12 category probabilities = 18 total values
X_holdout_combined_features[0]
```

340/340 0s 780us/step

```
[46]: array([ 1.41787788e-11,  1.74825016e-10,  5.77761927e-12,  4.99887562e-11,
            1.35361067e-09,  1.15742314e-08,  1.85001459e-09,  7.77609435e-11,
            4.40123042e-11,  1.00000000e+00,  7.29537430e-10,  3.35666583e-09,
           -4.12780575e-01, -4.23497271e-01, -1.52634330e-01,  1.65405516e-01,
           -4.24747638e-01, -4.37688430e-01])
```

Get predictions from model and print classification report and confusion matrix results

```
[47]: predictions_unseen_data = rfc_secondary.predict(X_holdout_combined_features)

print_results(y_holdout_data, predictions_unseen_data)
```

Classification report

	precision	recall	f1-score	support
0	0.89	0.94	0.91	186
1	0.83	0.86	0.84	203
2	0.99	0.98	0.98	2118
3	0.96	0.94	0.95	861
4	0.97	0.97	0.97	1570
5	0.85	0.84	0.84	417

6	0.95	0.94	0.94	431
7	0.93	0.94	0.94	982
8	0.91	0.91	0.91	382
9	0.98	0.98	0.98	1514
10	0.83	0.85	0.84	399
11	0.98	0.99	0.98	1804
accuracy			0.95	10867
macro avg	0.92	0.93	0.92	10867
weighted avg	0.95	0.95	0.95	10867

Confusion matrix

```
[[ 174   1   2   0   0   0   0   1   0   5   1   2]
 [   3 175   0   3   1   4   0   0   0   1 16   0]
 [   3   0 2078   1   6   0   1 21   0   2   0   6]
 [   0   6   5 807   5   3   1   4 12   3 12   3]
 [   0   1   7   7 1519   9   3   8   0   0   9   7]
 [   0   8   0   4   5 349   4   1 21   2 16   7]
 [   3   0   0   0   8   0 403   9   0   6   2   0]
 [   1   6 17   4   8   1   3 924   2   2   8   6]
 [   0   0   0 11   1 17   0   5 348   0   0   0]
 [ 11   3   0   3   3   1   5   5   0 1477   5   1]
 [   0 11   0   5   8 19   1   9   0   4 341   1]
 [   0   1   0   0   7 10   1   3   0   2   2 1778]]
```

[Parallel(n_jobs=8)]: Using backend ThreadingBackend with 8 concurrent workers.

[Parallel(n_jobs=8)]: Done 34 tasks | elapsed: 0.0s

[Parallel(n_jobs=8)]: Done 184 tasks | elapsed: 0.0s

[Parallel(n_jobs=8)]: Done 200 out of 200 | elapsed: 0.0s finished

It can be seen that the model performs just about the same on unseen data, which indicates that the model is generalizing well and is not overfitting to the training data.

[]: